Low-Frequency RoPE Subspace for Near–Far Sparse Attention

Anonymous authors

***Abstract*—Sparse attention must retain useful candidates at low selection cost. We study a near–far page-to-token cascade sharing 32 low-frequency RoPE scoring coordinates, original-dimensional attention on selected tokens, and absolute budgets: protected sink and window lengths, a near extent, explicit near page and token quotas, far capacity by subtraction. Ratio-coupled parameterizations instead silently reallocate far capacity: one expression links the far token budget to three separate knobs, and archived diagnostics exhibit all three resulting failures, including an exact integer threshold. Absolute quotas make both regions' capacities explicit and non-negative by construction. Author-reported 11-task RULER 32K AVGs are 58.30 and 59.29 under near quotas (16,640) and (8,320) on Qwen3-8B; a Qwen3-32B comparison reports 50.64 sparse versus 52.22 dense. These descriptive comparisons neither isolate mechanisms nor establish non-inferiority, and separate historical timings show slower sparse decoding. Establishing benefit for the compact 4/128-protection design still requires matched quality and preparation-inclusive complete-request cost.**

***Index Terms*—sparse attention, shared representation, regional retrieval, long-context inference**

I. Introduction

In dense long-context decoding, each new token attends to a growing history of cached keys and values. Sparse attention restricts final attention to selected tokens, but selecting them requires scoring, ranking and memory accesses. The design goal is to preserve task-relevant tokens within a fixed budget while keeping selection cheaper than the attention work it removes. Both candidate selection and final attention-set size are therefore central design problems.

Cascades such as AsyncTLS shortlist pages using inexpensive summaries before token refinement [4]. Mean pooling and sign-aware extrema, summaries our selector also uses, do not rank pages by their total attention weight. Following ClusterKV's token/page perspective [8], Figure 1(a) shows a single Qwen3 query where fewer missed mass-ranked reference pages coincide with more missed reference-page mass. Unequal regional allocations and realized token counts preclude a controlled comparison, but motivate examining how scoring and budgets jointly shape candidate access.

Candidate access also depends on how capacity is divided between recent and distant history. A global TopK can devote all available slots to one region, whereas regional quotas can reserve nominal capacity for each. In the diagnostic that selects whole bins in Figure 1(b), separate near and far reservations raise the mean retained fraction of far-region attention weight from .6759 to .7207, but reduce it in 43 of 128 correlated observations. The design must address both stages; its parameterization matters as much as its values, since ratio-derived quotas silently reallocate the opposite region's capacity, and admitting far pages alone does not reserve final slots for their tokens.

We study a two-level selector that couples a shared scoring view with absolute regional budgets. Both stages use a fixed set of 32 coordinates comprising 16 complete low-frequency RoPE pairs; near pages and near tokens receive explicit integer quotas, and far capacity follows by subtraction from the total cap. Among the three 32-coordinate subsets in Figure 1(c), low-frequency pairs retain .84 of full-dimensional reference mid-TopK mass on average, random coordinates .51, high-frequency pairs .24; the random subset does not preserve RoPE pairs. Selected tokens use original-dimensional Q/K/V in one final attention operation. Compact coordinates narrow scoring dot products, but ordinary token scoring still spans full history before masking; fewer shortlisted pages alone do not remove that work. We also formulate a prefill-informed far gate motivated by the layer variation in Figure 1(d); it is disabled by default and requires separate evaluation.

We contribute an absolute-budget two-level selector with shared low-frequency coordinates, a unified analysis of implicit budget coupling in ratio parameterizations, and separate evaluations of task quality, regional coverage and execution cost. Under the archived ratio interface, one expression in the page cap, near share and fill fraction determines the far token budget, so raising any of the three can zero it. On complete 11-task RULER32K sets, reported Qwen3-8B AVG is 0.99 points higher at near quotas (8,320) than at (16,640), with quotas changing together. A separate 11-task RULER32K comparison on Qwen3-32B shows a 1.58-point sparse deficit relative to dense attention. These descriptive comparisons do not isolate either mechanism, and the archived sparse configuration decodes more slowly than dense attention at every measured length. Establishing a quality–latency benefit for the compact 4/128-protection design requires matched task quality and complete-request measurements, including preparation, maintenance and selection.

II. Related Work

Quest supplies sign-aware page-extremum bounds for far-page scoring, SparQ approximates attention from query-selected key coordinates, and HISA prunes the attention set hierarchically [1], [2], [3]; SparQ protects a recent-token window, and HISA forces first and last blocks into its coarse candidate set. AsyncTLS already combines a page-to-token cascade with coordinate selection, quantization, per-query normalization followed by grouped-query-attention (GQA) probability aggregation, and final attention over original KV [4]. Rather than proposing a new cascade, we study these existing components with complete low-frequency RoPE pairs shared across both stages and absolute near–far budgets, distinguishing page-level eligibility from final token capacity.

The selector's own reads also matter. Token-Picker uses conservative probability estimates from partial key-bit chunks to stop further key-chunk and value transfers for pruned tokens [13]; our selector instead ranks candidates with a fixed coordinate subset under page and token budgets, keeping original-dimensional KV. TLI keeps both selection stages in bf16: AsyncTLS token keys are INT4, FFD scans 2-bit keys, and official HISA pools block summaries in FP8 [4], [5], [3]; our efficiency comes from candidacy structure, regional quotas and page-level narrowing, not numeric compression. KV compression and placement methods address complementary storage and movement constraints [6], [7], [8]; a smaller scoring view retains original-KV capacity without lowering complete-request latency.

III. Near–Far Selection in a Shared Subspace

To control both selection cost and candidate access, the cascade uses two explicit mechanisms: a shared scoring view and absolute regional budgets. One 32-coordinate map supports both L1 page screening and L2 token refinement, separating low-dimensional selection from original-dimensional attention (Fig. 2). The cascade

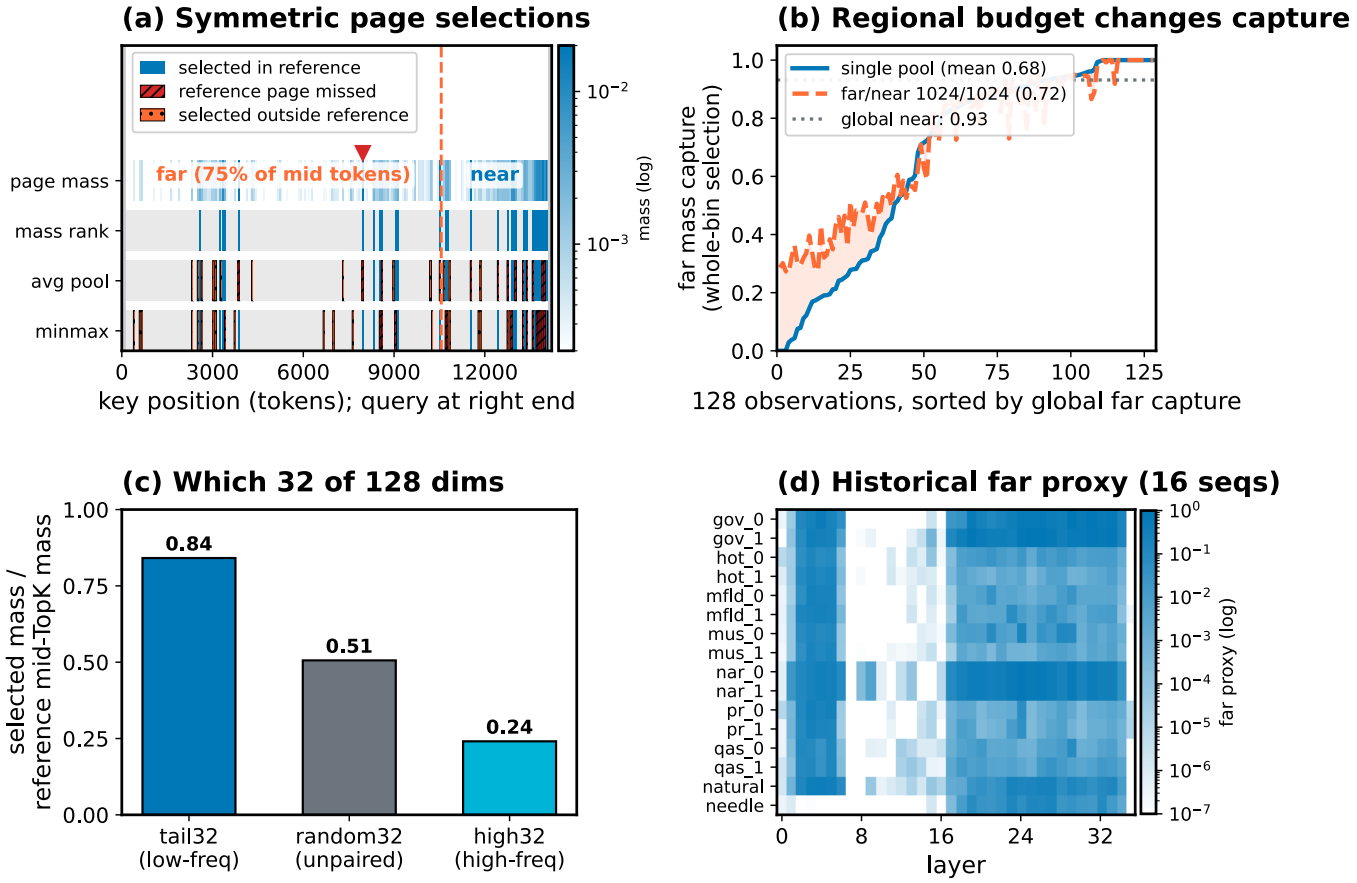


Fig. 1. Four motivating diagnostics; Qwen3-8B, 128/128 prefix/recent protection in (a)–(c). (a) MuSiQue sample 1, layer 24, query 14212. The top strip shows summed attention per page, not a selection. Below: stored mass-rank, avg and minmax selections, 32 pages each, using 1989/1989/2048 tokens under a 2048 cap. Avg/minmax miss 13/17 reference pages holding 59.65%/26.46% of reference-page mass. Blue: selected reference pages; red hatching: missed reference; orange dots: selected outside the reference. Mass-rank is a whole-page comparison reference, not an optimal oracle, exact token TopK or DSA. Avg and minmax share 32 coordinates and head aggregation; the reference uses full-dimensional attention weights. Different regional allocations and realized counts preclude inferring near/far specialization or task quality from these misses. (b) Four sequences contribute 128 correlated observations, ordered by global far retention. A global whole-bin cap of 2048 is compared with separate far/near caps of 1024/1024, without refill; red shading marks lower capture. Retention divides selected regional mass by full regional mass; the dotted line is global near retention. (c) Selected mass relative to full-dimensional reference mid-TopK mass; random32 does not preserve RoPE pairs. (d) One 16-by-36 historical single-query proxy matrix gives 24/36 layer medians below .01 and 12–32 low-proxy layers per sequence. It uses the last sampled query, group-logit summation before softmax and far interval $[64, t-2048)$ (S6), unlike the current gate. These are diagnostic selections and profiles, not task-quality or decode-safety results.

operates only during decoding; prefill remains dense. L1 assigns separate page quotas to the near and far regions, so pages in one region do not compete for the other region's allocation. Ordinary L2 scores the full history, then applies the retained-page mask and regional token selection; selected indices merge with protected positions for attention using the original Q, K and V, and recomputing retained logits does not recover omitted tokens.

A. *Shared low-frequency coordinates*

Both regions and both selection stages use the same coordinate map within each layer and KV head, while each query head is scored separately. For Qwen3-8B, we select post-RoPE coordinates $J = \{48, \dots, 63, 112, \dots, 127\}$ from each 128-dimensional query and key. The indices retain 16 complete low-frequency split-half RoPE pairs (Supplementary Fig. S1). With coordinate-selection matrix $E_J \in \mathbb{R}^{128\times 32}$, the scoring vectors are $z_q = E_J^\top q$ and $z_i = E_J^\top k_i$. The retained pairs' slow positional rotation motivates this selection, although it does not guarantee preservation of important tokens. The compact implementation uses width-32 scoring tensors, whereas the historical quality runs used 32 active coordinates in width-128 L2 tensors (S7).

L1 aggregates selected key coordinates into page summaries for coarse ranking; L2 refines candidates with per-token scores in the same coordinate view. The two scoring views occupy separate buffers. Original KV remain full width. Compact scoring narrows each dot product fourfold, but view preparation and maintenance add work without compressing the original KV cache (S7).

B. *Absolute budgets and regional eligibility*

In decoding, t indexes the current token from zero across the full sequence including prefill. Let s and w denote the protected sink and recent-window lengths, with protected set $\mathcal{P}_t = \{i : 0 \le i \le t,\ i < s \text{ or } i > t - w\}$ and unprotected middle $\mathcal{M}_t = \{0, \dots, t\} \setminus \mathcal{P}_t$; overlapping protection counts once. Near is a suffix of $\mathcal{M}_t$ with explicit page-aligned extent L_n; far is the rest. Near excludes the protected recent window.

Regional page caps and token quotas constrain the cascade (Supplementary Fig. S2). Near pages use query–mean-key dot products to summarize page-level relevance; far pages use sign-aware extrema to

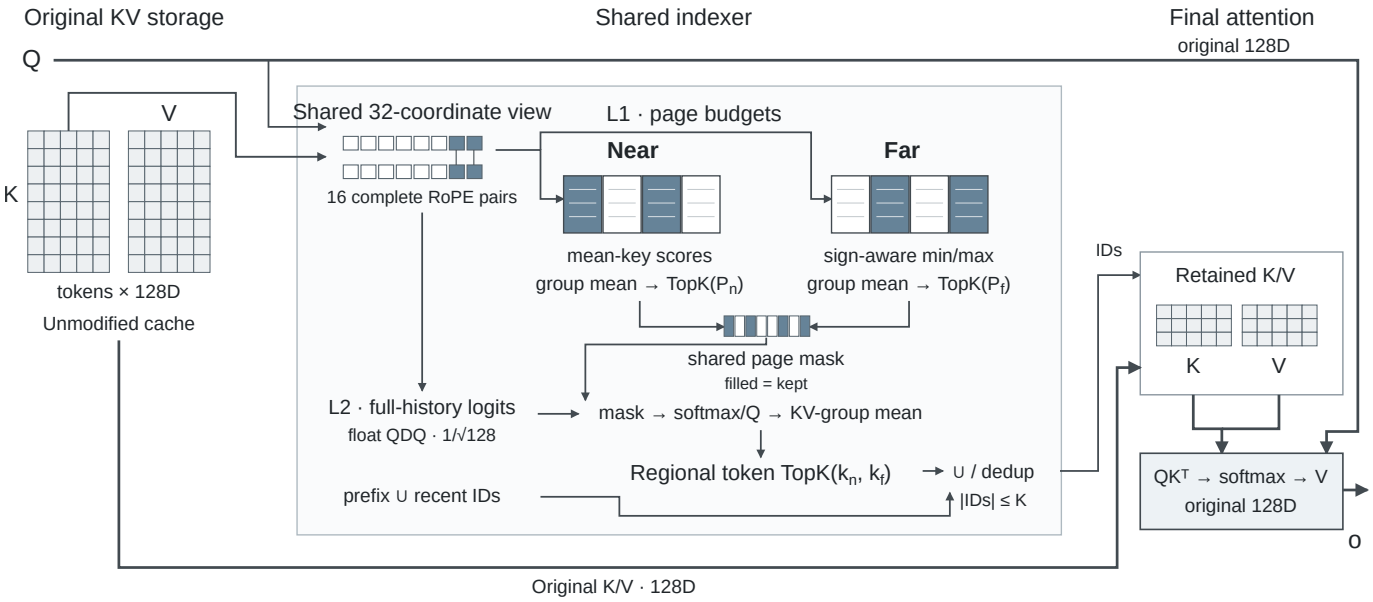


Fig. 2. Architecture of shared-subspace selection, with the optional gate inactive. A shared coordinate view feeds regional L1 page scoring and ordinary full-history L2 scoring. Page masks limit eligibility; per-query softmax and KV-group probability averaging precede regional token selection. Protected positions join deduplicated IDs within K. Near excludes the protected recent window. IDs define support over original-width K/V; physical gathering is backend-dependent. Original Q follows a separate path. Arrays are schematic.

bound shared-space token scores, a heuristic assignment. Summaries cover full physical pages including protected boundary positions; L2 separately limits token regions (S7). For a nonempty page b, let $m_{b,j} = \min_{i \in b} z_{i,j}$ and $M_{b,j} = \max_{i \in b} z_{i,j}$. Quest's construction [1] gives

$$U_b(z_q) = \sum_j (z_{q,j}^+ M_{b,j} + z_{q,j}^- m_{b,j}), \tag{1}$$

where $z^+ = \max(z, 0)$ and $z^- = \min(z, 0)$. Each signed coordinate contribution is bounded by its chosen extremum, so $z_q^\top z_i \le U_b(z_q)$ for every shared-space key in b. This does not bound original-dimensional logits. Extrema may come from different tokens, making the bound loose enough to misrank pages and miss important tokens. Average-key scores lack this per-token bound.

For total cap $K \ge |\mathcal{P}_t|$, define the middle-token budget $K_{\text{mid}} = K - |\mathcal{P}_t|$. The interface exposes absolute quantities: near page quota n_p, near token quota n_k with $0 \le n_k \le K_{\text{mid}}$, and far capacity by subtraction, $K_1 - n_p$ far pages and $K_{\text{mid}} - n_k$ far tokens. Each quota is explicit; neither region's capacity can change without changing its own parameter. With enough positive-probability candidates, L2 cannot recover unprotected tokens from L1-rejected pages. S7 details scarce-support and masked zero-tie exceptions.

Ratio-coupled interfaces hide this subtraction. The archived parameterization derives $n_p = \max(1, \text{round}(\beta K_1))$ and $n_k = \min(\lfloor \gamma n_p B \rfloor, K_{\text{mid}})$ with ties-to-even rounding, so the far token budget reduces to

$$k_f = \max(0,\ K_{\text{mid}} - \min(\gamma n_p B,\ K_{\text{mid}})). \tag{2}$$

Three knobs, K_1, β and γ, therefore reallocate far capacity through one expression, while far pages remain admitted at $k_f = 0$: page eligibility and token capacity are separate constraints. Archived diagnostics exhibit all three reallocations, including an exact integer threshold at which raising K_1 zeroes far tokens (Q2, S6). Explicit quotas are immune by construction: n_k is bounded by K_{mid} and far capacity is stated, not derived.

With disjoint protection and sufficient candidates, the compact configuration ($K = 1024$, $s = 4$, $w = 128$) leaves 892 middle slots per KV head. Near quotas $(n_p, n_k) = (8, 320)$ leave 572 far slots, the $(16, 640)$ alternative 252 (Supplementary Fig. S2). The partial prefix page can reselect up to four protected tokens, so far coverage is nominal.

C. GQA selection and original-dimensional attention

GQA assigns four of Qwen3's 32 query heads to each of its eight KV heads. L1 averages their page scores before regional TopK. Ordinary L2 computes each query's full-history logits, applies the shared page mask, then softmax. Group-averaged probabilities determine regional token TopK. L2 uses scale $1/\sqrt{128}$ before softmax even with 32 scoring coordinates. Because softmax precedes group averaging, this rule generally differs from applying softmax to group-averaged logits. A smaller L1 shortlist restricts eligibility without reducing ordinary L2's full-history dot products.

For KV group h, let $\mathcal{S}_{t,h}$ be the deduplicated union of protected positions and regional winners, with $|\mathcal{S}_{t,h}| \le K$. Each query g uses this support for its own full-dimensional attention ($d = 128$):

$$o_{h,g} = \sum_{i \in \mathcal{S}_{t,h}} \frac{e^{q_{h,g}^\top k_{i,h}/\sqrt{d}}}{\sum_{j \in \mathcal{S}_{t,h}} e^{q_{h,g}^\top k_{j,h}/\sqrt{d}}} v_{i,h}. \tag{3}$$

Each query applies one softmax over the merged indices, avoiding the addition of separately normalized regional outputs [9]. Full-dimensional retained logits still use a restricted denominator, which can change the output. These indices define a shared logical support. Physical KV gathering and reuse depend on backend and length; short sequences can use a dense masked path.

D. Prefill-informed far gating

The optional gate, disabled by default, exploits layer variation: layers whose mean prefill far-attention proxy falls below ϵ skip far selection during decoding, and the first decode-path call freezes the

TABLE I
Compact design parameters and execution policy. These settings define the proposed implementation, not the historical runs. K includes unique protected tokens.

Parameter	Setting
Final token cap K	1024
Page size; L1 page cap K_1	64; 128
Near page quota n_p (far)	8 (120)
Near token quota n_k (far)	320 (572)
Near extent L_n	quarter of middle, page-aligned
L1/L2 active coordinates	32/32
Prefix / recent protection	4/128
Prefill; decode layers	dense; all sparse

decision until input reset (Supplementary Fig. S3; S5). OFF skips far L1/L2 scoring and selection in the supported fixed-quota path, refilling near L2 within unchanged near pages; removing far support can still change softmax denominators and group-averaged near rankings. Calibration uses a token-defined region that can differ from the page-aligned skipped region, the threshold does not bound decoding error, and preparation, sequence-wide buffers and normalization remain (S5). Gate quality and total cost require separate evaluation.

IV. Experimental Setting

We evaluate three aspects separately: task quality (Q1), regional candidate coverage (Q2) and execution cost (Q3). Table I defines the compact Qwen3-8B design: dense prefill, sparse decoding and a 1024-token cap including four prefix and 128 recent tokens (the 4/128-protection design). Archived runs use ratios (α, β, γ) mapping to the same geometry via $n_p = \max(1, \mathrm{round}(\beta K_1))$, $n_k = \min(\lfloor \gamma n_p B \rfloor, K_{\mathrm{mid}})$ and page-aligned $L_n = \lfloor \alpha |\mathcal{M}_t| \rfloor$; $(8, 320)$ equals the archived $(.25, .0625, .625)$ cell (S7). Archived quality runs applied 16-level quantize–dequantize to L2 token keys; we treat it as an implementation detail outside the mainline (S7). The historical five-task LongBench shared-subspace arm uses an earlier implementation: 128/128 protection and 32 active coordinates in zero-filled width-128 L2 keys. The labels "full-dimensional page/token scoring" denote reported variants without complete execution configurations; we retain them as nominal, not as controlled single-stage ablations. Synthetic selector and model-level decode timings come from separate studies without matched quality or cost; S6/S7 detail configurations and measurements.

Workloads and comparisons. Reported quality covers five-task LongBench [10] and 11-task RULER [11] (Qwen3-8B: 32K/64K, historical 128K; Qwen3-32B: 32K), plus ported Quest/AsyncTLS/JSSA baselines and a Llama-3.1-8B arm at 32K. Three nominal sparse LongBench variants lack verified factorial isolation. FullKV is dense and uncapped (RULER/decode timing). Native protocols retain disclosed prefill/decode; policy-matched adaptations need separate reporting. Budget matching aligns final unique-token caps K including protection; realized counts and head sharing affect KV traffic, so equal K alone establishes neither matching (S7).

Metrics. Primary quality is task-score AVG over each fixed benchmark/length set; per-task gains and losses explain AVG without a zero-loss requirement. Policy-matched controls require matched models, samples, lengths, scorers, phase policies and aligned sparse-token budgets; incomplete sets lack full-set AVG. Regional retention (Fig. 1(b)) divides selected by full regional mass; panel (c) uses full-dimensional reference mid-TopK mass. Neither selector time (synthetic single-layer) nor synchronized TPOT (model-level intervals after the first output token) measures preparation-inclusive complete-request latency. Unpaired quality and timing studies do not define a quality–latency frontier; AVG differences are descriptive, not statistical non-inferiority evidence.

V. Evaluation and Discussion

A. Q1: Does reduced scoring width preserve task quality?

On HotpotQA, Qasper, MuSiQue, GovReport and RepoBench-P [10], five-task means sit within a descriptive 0.33-point spread: shared-subspace 47.588, nominal full-dimensional page 47.910 and token 47.666 (Supplementary Fig. S4); the shared-subspace arm sits 0.322 and 0.078 lower. These differences describe recorded configurations, not evidence that reduced width preserves quality: the variants are not verified single-factor interventions, and the shared-subspace run uses earlier 128/128 protection. All three arms are sparse; dense FullKV is a different reference.

B. Q2: How do regional budgets constrain candidate access?

Regional quotas lower far retention in 43/128 correlated observations despite improving its mean (Fig. 1(b)); this is not task-quality evidence. Reported 8B 32K AVG rises 58.30→59.29 when near quotas move $(16, 640) \to (8, 320)$ at $K_1 = 128$, $K = 1024$, protection 4/128 (archived $\beta = .125 \to .0625$); far pages rise 112→120 and the far token budget 252→572 jointly. Halving the near page quota to 4 (archived $\beta = .03125$) scores 58.07 at $K_1 = 128$ with far budget 732, and 58.96 for the $K_1 = 256$ variant with unchanged quotas; both sit below the tested $(8, 320)$, with CWE up 3.9 and FWE down 9.67. Data/scorer records match these sparse arms; predictions remain unreplayed.

Implicit coupling failures. Archived diagnostics exhibit all three reallocations of Eq. (2). Integer enumeration of the production formulas locates the first far-zero page cap at $K_1 = 156$ for $K = 1024$ ($B = 64$, $\beta = .125$, $\gamma = .625$, 128/128 protection); far budget 8 remains at $K_1 = 155$, so capacity appears to grow while far tokens vanish. Archived task scans collapse HotpotQA from 53.6–57.3 to 37.7 exactly on zero-far-budget arms, graded by far fraction (16.4 points at .875, about 1 at .125). At the production cell $(\alpha, \beta) = (.25, .125)$, free competition ties the best γ within .001 in retained-mass diagnostics (.852 versus .8515), while $\gamma = 0$ reads as far-free but zeroes the near token quota, dropping retained mass to .827. The historical 128K RULER arm computed $k_f = 0$ under its reported budgets (S6). These archived diagnostics unify the three anomalies but are not new controlled experiments.

Reported 32K AVG gaps to dense FullKV are -0.09 (8B, different scorers) and -1.58 (32B, matching manifests), across machines. At 64K, 8B AVGs are 48.85/48.05 for near quotas (8,320)/(16,640) versus 48.54. The sparse arms share a machine and scorer hashes but differ in prediction scripts; the dense anchor differs in scorer and legacy-partial identity, and missing paired intervals preclude causal conclusions (S8).

Ported baselines. Author-reported AVGs (nominal 512): 56.00 (ours) and 48.33 (AsyncTLS), a reported 7.67-point gap; the new AsyncTLS/FASA artifacts sit outside the repository snapshot, and unique-token/protection/phase matching is unverified. At $K = 1024$, 59.29 compares with JSSA's 58.7 (nominal 1024), Quest's 42.85 and FASA's 18.64; $K = 384$ scores 15.00. Cross-machine Llama-3.1-8B: 83.88/86.19 at near quotas (8,320)/(16,640). S8 gives provenance, CWE's 43.4-point loss and absent paired intervals.

At 128K, the (8,320)/(16,640) arms score 45.73/46.02 on one machine; 45.73 is 98.49% of the separate-run dense anchor, 46.43. Historical RULER 128K yields 23.98 (sparse) under budgets implying $k_f = 0$, execution unverified, the 22.45-point loss unexplained (S6).

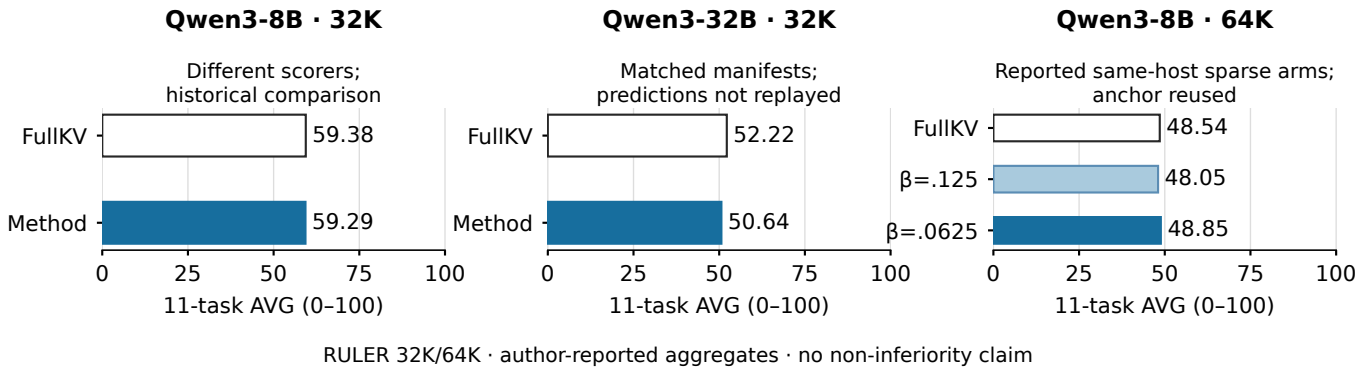


Fig. 3. Author-reported RULER 32K/64K AVG: 11 tasks, 100 examples/task/arm. FullKV is dense; sparse arms report $K = 1024$, protection 4/128. At 32K, near quotas (8,320) ($\beta = .0625$, 8B) or (16,640) ($\beta = .125$, 32B); at 64K, both are shown for 8B. The 32K comparisons span machines (8B scorers differ; 32B manifests match, predictions unreplayed); the 64K sparse arms are reported on one machine, the FullKV anchor a separate run. Gate status and scoring dimensions are unverified. These descriptive scores establish neither non-inferiority nor matched quality–cost benefit (S8).

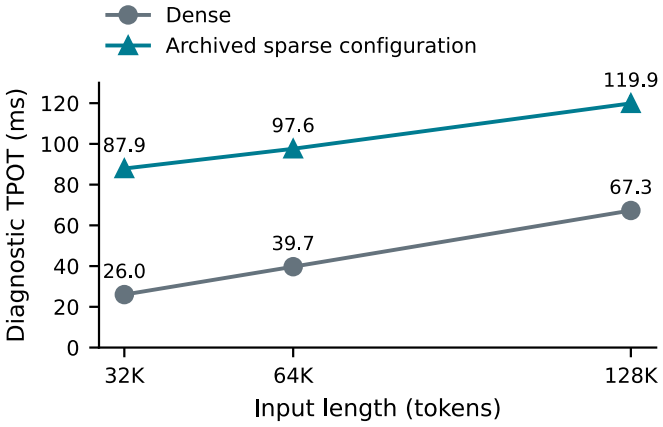


Fig. 4. Author-reported model-level decode TPOT for an archived sparse configuration and dense FullKV, Qwen3-8B with dynamic cache. Fixed 128-token output gives 127 after-first intervals, including synchronization and argmax. The archived $(.125, .375, .625)$ setting differs from the selected 4/128 candidate; these medians measure no complete-request E2E and establish no matched quality–cost comparison (S6).

Five-task LongBench means are 26.09 (protection-only, 256 tokens) versus 47.91 (reused sparse, 1024); retrieval and budget both change.

C. Q3: Does compact scoring reduce execution cost?

Model-level decoding. The reported sparse configuration is slower than dense FullKV at every measured length (Fig. 4). Sparse TPOT medians are 87.9, 97.6 and 119.9 ms at 32K, 64K and 128K, compared with dense medians of 26.0, 39.7 and 67.3 ms. The relative penalty shrinks with length, but absolute sparse TPOT increases: no decode-time benefit. They include synchronization and argmax, use a non-compact configuration, and cannot locate causes or predict the compact implementation (S6).

Local selector cost. A separate two-round synthetic experiment (single layer, batch one, one query, 64K QKV) reports steady-state summaries of 1.3502 ms (full-dimensional page variant), 1.3598 ms (shared-subspace) and 1.4634 ms (token variant); shared-subspace runs 0.0237 ms slower in round one and 0.0046 ms faster in round two. Missing raw sequences and an unresolved sampling discrepancy prevent assessing these small differences (S6). These local timings identify no bottleneck, lack quality pairing and measure no complete-request cost.

D. Optional gate: prefill is not a decode guarantee

Low prefill attention does not establish safe decode-time skipping. An archived input–layer record pairs a prefill far-attention proxy below .05 with a decode-tail mean of .1678 over five queries: a proxy shift, not measured quality loss. Gate thresholds .05 and .10 yield five-task AVGs of 47.876 and 47.692 against the reused 47.910 reference, down .034 and .218, with three tasks declining per arm, led by HotpotQA's 1.05 and 1.26. AVG remains primary; no per-task zero-loss requirement applies. Neither these changes nor zero-containing task-bootstrap intervals establish non-inferiority; missing scoring configurations, paired predictions and matched costs leave gated quality and net savings unverified (S6).

Design implications and limitations. Two constraints recur: reduced scoring width narrows each dot product while ordinary L2 still scores the full history, and far-page quotas admit candidates without guaranteeing far-token capacity. A useful operating point reduces complete-request latency at matched quality, counting preparation, maintenance and selection once. The archived sparse configuration decodes slower than dense, selector timings establish no request-level benefit, and matched measurements for the compact design remain missing; historical gate summaries lack complete configuration identity.

VI. Conclusion

We studied shared low-frequency page/token scoring with absolute regional budgets and original-dimensional attention on selected tokens. The analysis separates narrower scoring from reduced work, far-page reservation from far-token coverage, and explicit quotas from ratio-derived capacity, whose coupling expression permits all three archived reallocations. Qwen3-8B RULER 32K AVG is 0.99 points higher at near quotas (8,320) than at (16,640); the near-page-4 arms score 58.07/58.96, below (8,320). A Qwen3-32B sparse arm trails dense by 1.58 points; at 128K (8,320) reaches 98.49% of a separate-run dense anchor, and Llama-3.1-8B's (8,320) trails (16,640) by 2.31 points at 32K across machines. These results neither isolate mechanisms nor establish non-inferiority; the 128K loss and slower sparse decoding remain configuration limitations, and the compact 4/128 design still requires matched quality and complete-request cost.

Content Generated by AI. OpenAI models via Codex assisted drafting, revision, translation, figures, code and source checks; no model-version identifiers were exposed.

References

[1] J. Tang, Y. Zhao, K. Zhu, G. Xiao, B. Kasikci, and S. Han, "Quest: Query-Aware Sparsity for Efficient Long-Context LLM Inference," in *ICML*, 2024. https://arxiv.org/abs/2406.10774.

[2] L. Ribar, I. Chelombiev, L. Hudlass-Galley, C. Blake, C. Luschi, and D. Orr, "SparQ Attention: Bandwidth-Efficient LLM Inference," in *ICML*, 2024, pp. 42558–42583. https://proceedings.mlr.press/v235/ribar24a.html.

[3] Y. Xu *et al.*, "HISA: Efficient Hierarchical Indexing for Fine-Grained Sparse Attention," in *COLM*, 2026. https://arxiv.org/abs/2603.28458v4.

[4] Hu *et al.*, "AsyncTLS: Efficient Generative LLM Inference with Asynchronous Two-level Sparse Attention," arXiv:2604.07815v1, Apr. 9, 2026. https://arxiv.org/abs/2604.07815v1.

[5] Z. Liu, Z. Ning, R. Li, X. Liu, Y. Song, M. Zhang, Z. He, and X. Qiu, "Faster Than Flash: Exploiting Attention Sparsity for Efficient Long-Context Decoding," in *ICML*, 2026, pp. 77913–77928. https://proceedings.mlr.press/v306/liu26da.html.

[6] H. Sun *et al.*, "ShadowKV: KV Cache in Shadows for High-Throughput Long-Context LLM Inference," in *ICML*, 2025, pp. 57355–57373. https://proceedings.mlr.press/v267/sun25b.html.

[7] Z. Wang et al., "MILLION: MasterIng Long-Context LLM Inference Via Outlier-Immunized KV Product QuaNtization," in *2025 62nd ACM/IEEE Design Automation Conference (DAC)*, 2025, pp. 1–7, doi: 10.1109/DAC63849.2025.11132862.

[8] G. Liu, C. Li, J. Zhao, C. Zhang, and M. Guo, "ClusterKV: Manipulating LLM KV Cache in Semantic Space for Recallable Compression," in *DAC*, 2025. https://doi.org/10.1109/DAC63849.2025.11132479.

[9] G. Shen, J. Zhao, Q. Chen, J. Leng, C. Li, and M. Guo, "SALO: An Efficient Spatial Accelerator Enabling Hybrid Sparse Attention Mechanisms for Long Sequences," in *DAC*, 2022, pp. 571–576. https://doi.org/10.1145/3489517.3530504.

[10] Y. Bai *et al.*, "LongBench: A Bilingual, Multitask Benchmark for Long Context Understanding," in *ACL*, 2024, pp. 3119–3137. https://aclanthology.org/2024.acl-long.172/.

[11] C.-P. Hsieh *et al.*, "RULER: What's the Real Context Size of Your Long-Context Language Models?," in *COLM*, 2024. https://arxiv.org/abs/2404.06654.

[12] G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis, "Efficient Streaming Language Models with Attention Sinks," in *ICLR*, 2024. https://arxiv.org/abs/2309.17453v4.

[13] J. Park, M. Kang, Y. Han, Y.-G. Kim, J. Shin, and L.-S. Kim, "TokenPicker: Accelerating Attention in Text Generation with Minimized Memory Transfer via Probability Estimation," in *DAC*, 2024, doi: 10.1145/3649329.3655953.