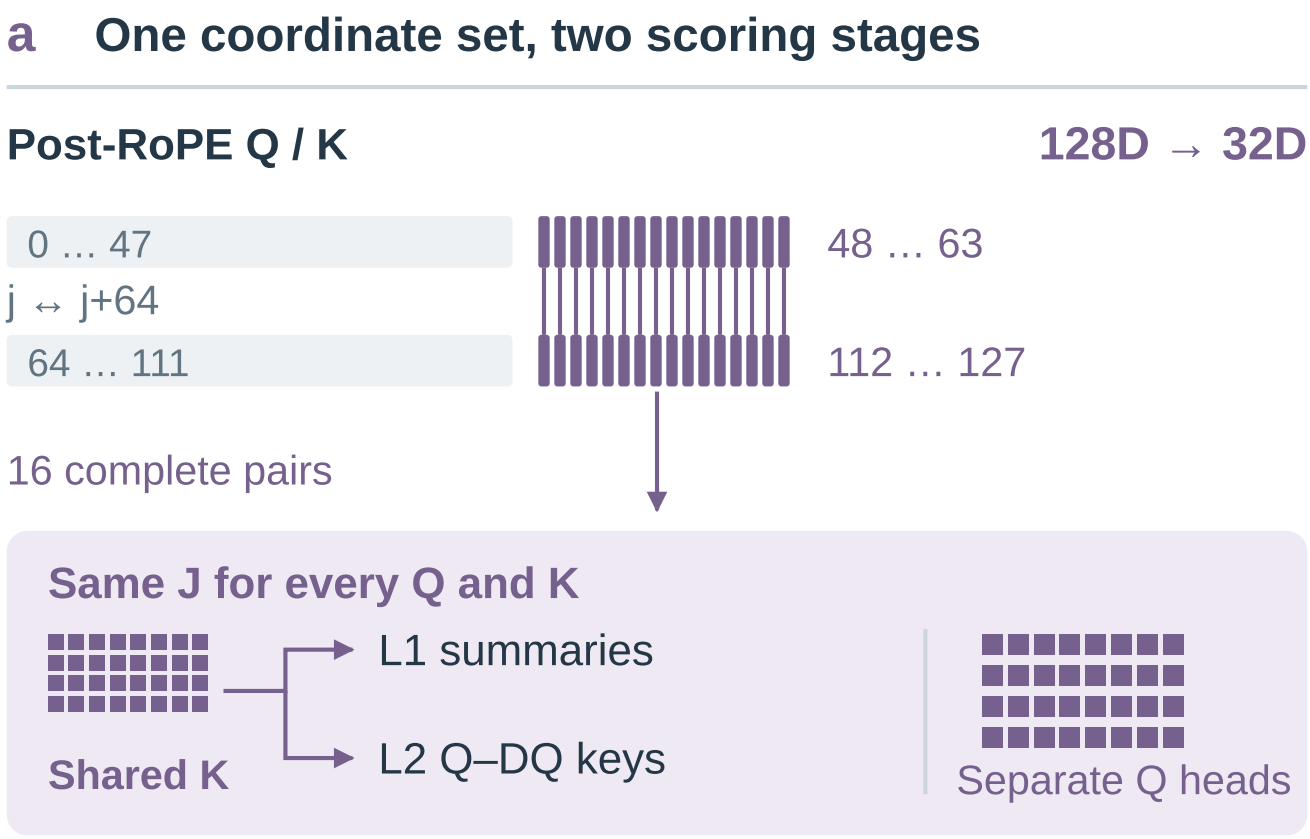
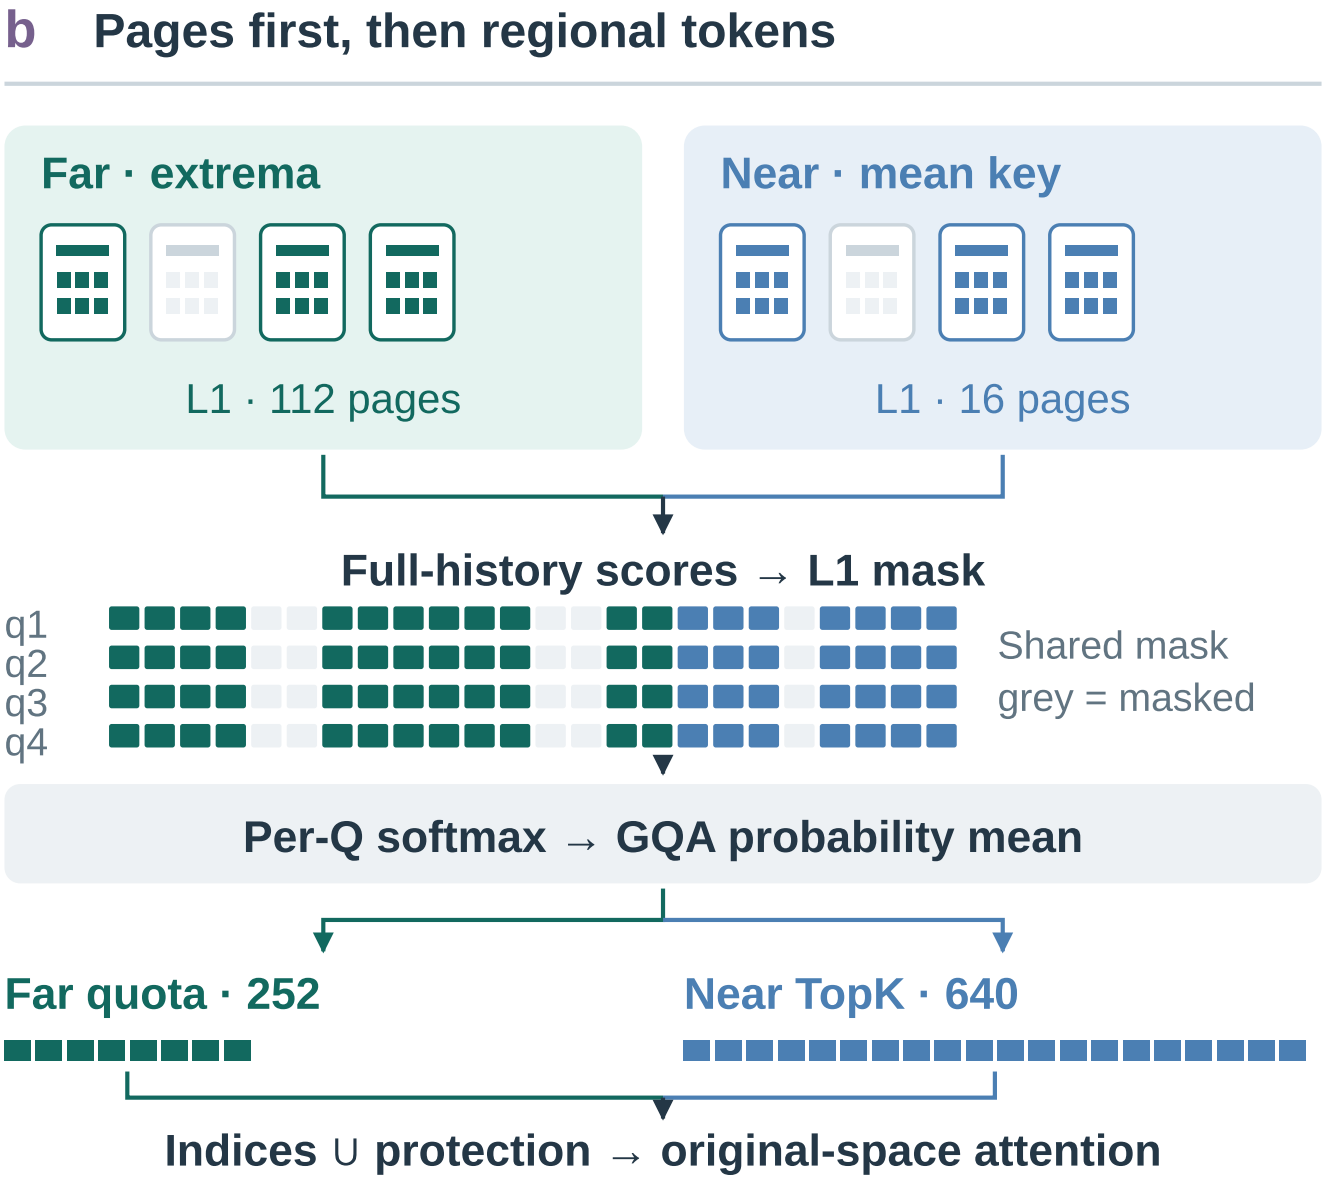
Supplementary Figure S1: Shared scoring coordinates

Detailed views of subspace selection, regional allocation, and the gate lifecycle under the assumptions stated in the main text.



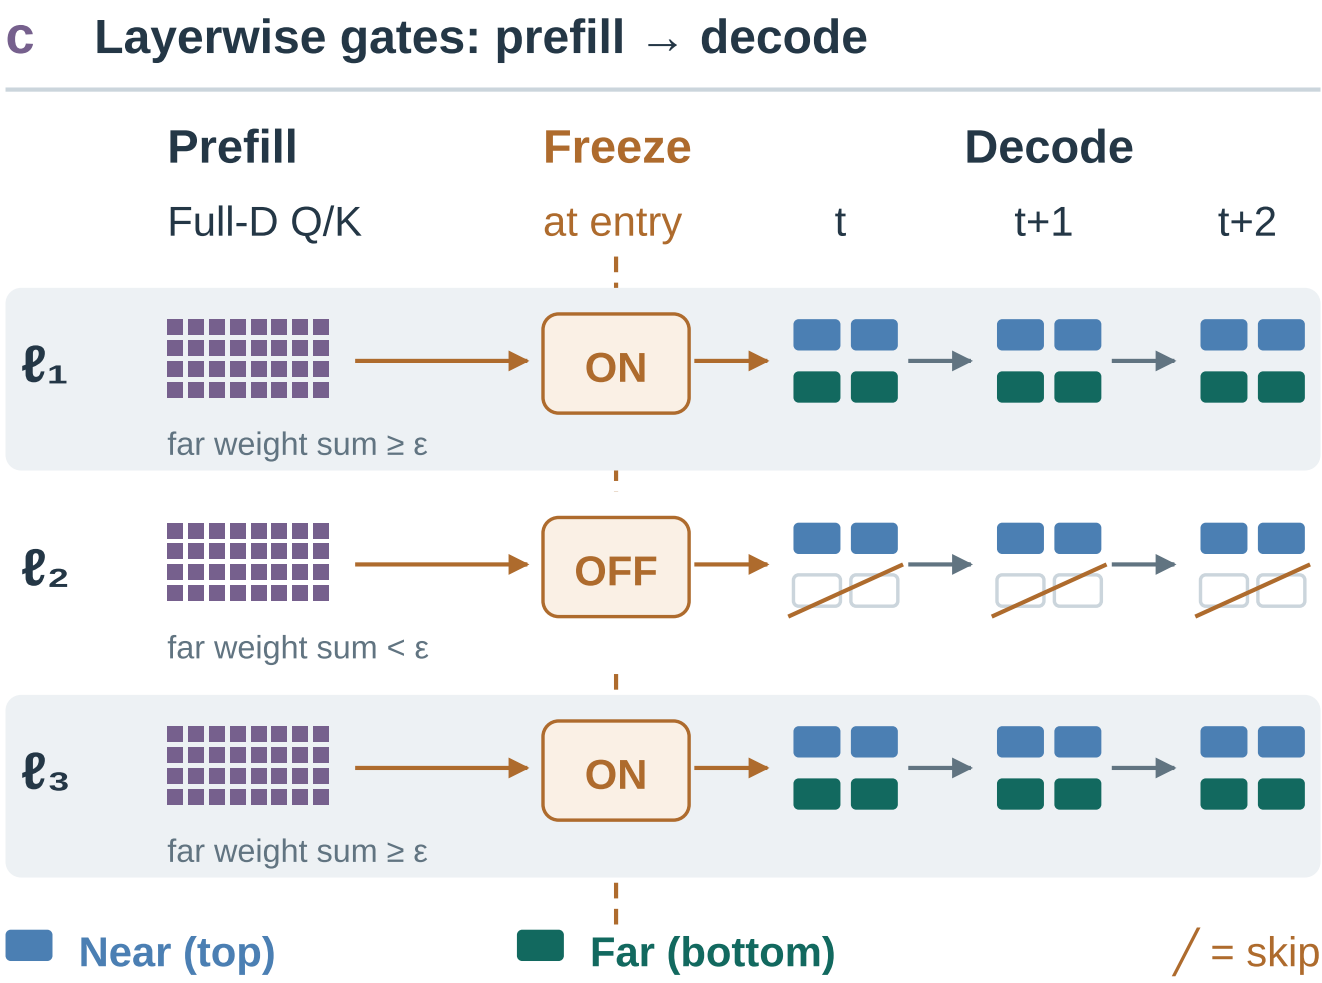
Subspace selection in the current compact implementation. The same 16 complete RoPE pairs feed L1 and L2 in both regions. The implementation stores floating-point quantize–dequantize values, not packed INT4. Selected IDs retrieve original 128-dimensional KV.

Supplementary Figure S2: Regional page and token allocation



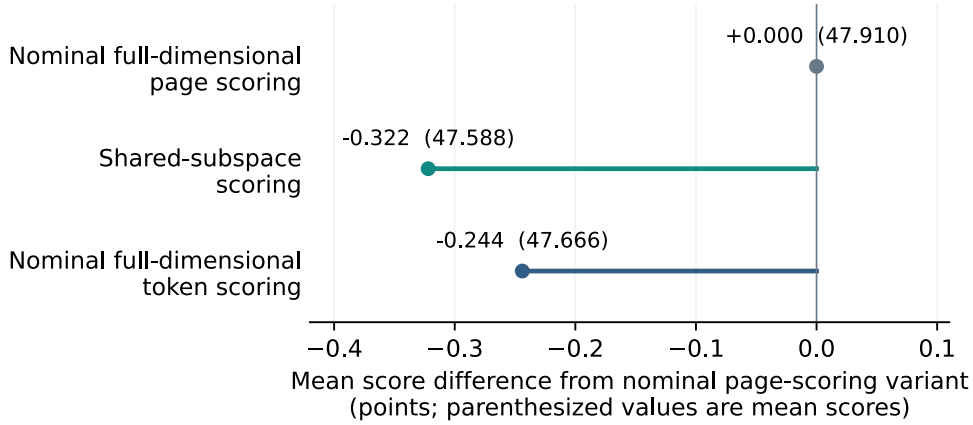
Regional selection and protected-token accounting. Under sufficient eligible capacity, $K_1 = 128$ gives 16 near and 112 far pages; the total 1024-token cap includes 132 protected tokens and nominal 640/252 near/far slots for the selected 4/128 configuration. Up to four prefix positions can be reselected from their partial page, reducing unprotected far coverage. The diagram's blocks are schematic. Per-query-head normalization precedes KV-group averaging.

Supplementary Figure S3: Layerwise prefill-to-decode gating



Layerwise gate state transitions and conditional execution. Prefill stores a decision; the first decode-path call freezes it. In the fixed numeric-quota path with valid length and nonempty far history, OFF skips far L1/L2 dot products and far selection. Only unused L2 quota refills eligible near tokens; the near L1 cap stays 16 pages. Global preparation and S-wide buffers/softmax remain. Layer states and repeated decode steps are schematic, not measured layer outcomes. The gate is disabled by default and has no matched quality–cost validation for the shared-subspace configuration.

Supplementary Figure S4: Descriptive configuration comparison



The five-task means are 47.910 (nominal full-dimensional page-scoring variant), 47.588 (shared-subspace configuration) and 47.666 (nominal full-dimensional token-scoring variant). The differences of 0.322 and 0.078 points from the shared-subspace configuration are descriptive. The nominal full-dimensional page-scoring variant lacks experimental metadata; the nominal full-dimensional token-scoring variant's L2-mode flag is not bound to an implementation revision that implements that mode. These labels therefore do not establish controlled stage interventions or a completed three-configuration ablation. Supplementary Note S7 gives the configuration evidence. None is dense FullKV, and no document-paired non-inferiority or matched quality–cost conclusion follows.

Supplementary Note S5: Gate observation and execution scope

When the optional gate feature is disabled, ordinary far selection remains active; OFF denotes a calibrated per-layer skip decision only when the feature is enabled. The optional far-region gate observes a layer's prefill queries, stores a pending far-skip decision, freezes it at the first decode-path call, and reuses it until a new input clears the state (Supplementary Fig. S3). That call can still consume question tokens, so it need not see the complete question. Calibration uses at most the final 256 available query rows, full-dimensional scores and per-query-head softmax; its averaged far-attention weight sum is compared strictly with the threshold. The threshold controls a prefill observation; it does not bound far-attention weight sum or task error during later decoding.

In the supported fixed-quota implementation (batch one, valid length and nonempty far history), OFF skips far L1/L2 dot products and selection. Unused L2 quota refills eligible near candidates without enlarging near's L1 page budget. Far-on, off-refill and off-no-refill are distinct interventions: removing far can change per-query softmax denominators and the aggregated near ranking. Consequently refill is not an isolated saving with an invariant near set. Preparation, quantization, repeat-to-float and S-wide buffers/softmax remain in the cost. The implementation does not support combining dynamic and static layer skipping or applying the dynamic gate to Llama. Batched state isolation, runtime state transitions and distribution-shift-triggered fallback have not been validated. Runtime adaptation of α, β, γ is outside the evaluated configuration.

The main-text gate is a candidate extension to the shared-subspace method. The archived gate task-level summaries do not identify generation-code revision, complete subspace flags or paired predictions, so they cannot validate this configuration. These implementation details qualify the mechanism rather than provide measured speed or quality gains.

Supplementary Note S6: Scope of the archived diagnostics

The following records retain the numerical detail behind the main-text limitations. They are author-archived diagnostics, not new experiments or matched validation of the shared-subspace method.

Motivating diagnostics. Figure 1 replaces the earlier over-budget whole-bin heuristic with the checked implementation from commit `8c4ff13`. A bin spans $[q-d_{hi}, q-d_{lo}+1)$ and is selected only if its full valid-token count fits the remaining 2048-token middle budget. Realized counts range from 257 to 2048. Near/far attention retention uses the same global selection: selected weight in each region divided by that region's total weight. Probabilities are softmax-normalized per query head, averaged within each KV group, then equally across eight KV groups; neither is a $K = 1024$ production result. Density uses bin mass divided by its exact valid-token count, without reconstructing within-bin token probabilities. The displayed query is chosen only for illustration: among queries beyond position 9000, its far share is nearest the pooled median; this is not a preregistered selection rule. Its 256 protected tokens have full-reference union mass 0.6304. Four sequences contribute two layers and 128 correlated query/layer observations; no independent-document uncertainty is estimated. Q/K snapshot and stripped O3 sidecar hashes are recorded, but their host-side tensors were not recovered. The coordinate diagnostic below averages eight task/sample/layer cells equally; its random coordinates do not preserve RoPE pairs. It neither isolates frequency nor measures production regional scoring or task quality.

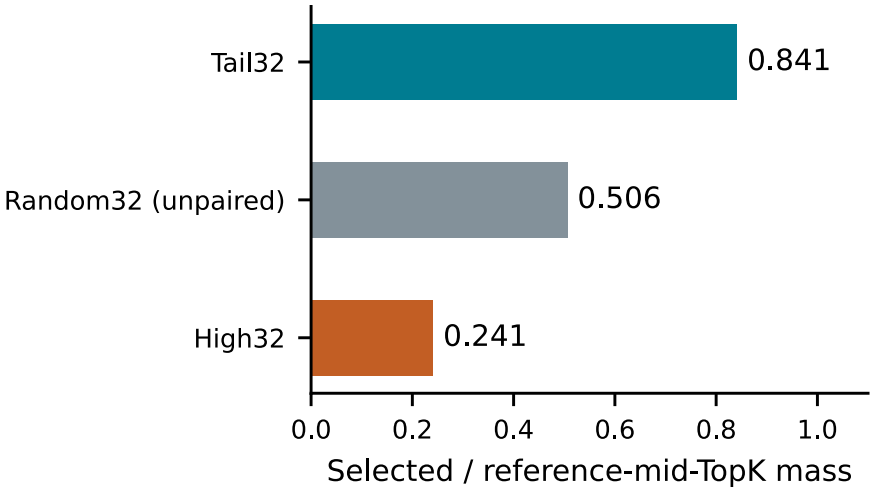


Figure 1: Secondary archived coordinate diagnostic: selected mass divided by full-dimensional reference mid-TopK mass, with middle cap 768 and separate 128/128 protection. The diagnostic page-average selector differs from production near-average/far-minmax scoring. Random32 is unpaired. No paired-random quality or output-error test is claimed.

Timing provenance. Figure 3 uses the archived sparse configuration $(\alpha, \beta, \gamma) = (.125, .375, .625)$, nominal $K_1 = 128$ and $K = 1024$, whose protection is not explicitly bound. Qwen3-8B bf16 runs use an NVIDIA H20, Transformers 4.56.2 and Torch 2.8.0 as reported by the source. The minimal derived plotting input retains these fields and the original repository SHA-256, while omitting host addresses and private paths; its hash is distinct from the original. E152v2 supplies author-reported medians, not raw repetition records. An author audit lists three dynamic records per arm and length with three measurements per record; its written interleaving order does not justify labeling every arm ABBA$\times$3. The current collector labels the backend route unverified. The 82.0-ms StaticCache point belongs to a different repair and repetition identity and is omitted. StaticCache construction occurs before the request timer. The E127 5.47-fold figure compares old and optimized TLI stage-time summaries, not dense FullKV. Neither it nor the E152 aggregates establishes a launch-bound cause. Source-data, collector and renderer hashes are supplied in `sources/evidence_bindings.json`.

Local selector timing and aggregation. The archived performance record reports two GPU-crossed rounds for a later implementation. Its associated source has a compact shared-subspace L2 path, distinct from the earlier LongBench implementation in Note S7; source inspection alone does not certify the timed execution. The per-round values are 1.3635/1.337 ms (nominal full-dimensional page-scoring variant), 1.3872/1.3324 ms (shared-subspace configuration), and 1.4794/1.4473 ms (nominal full-dimensional token-scoring variant). Its reported two-round medians are 1.3502, 1.3598 and 1.4634 ms, respectively. Midpoint medians computed from each pair are 1.35025, 1.35980 and 1.46335 ms; these are arithmetic recomputations, not additional measurements. The record does not supply per-stage medians or raw sequences and therefore does not establish the earlier description as a sum of five stage medians. The synthetic single-layer, batch-one, single-query 64K scope is distinct from complete layer or request execution and is not paired with the quality runs. The record describes 60 steps, 40 warmups and 20 steady samples, whereas the previously inspected harness implies 100 total and 60 steady samples. This discrepancy remains unresolved. The reported

shared-subspace timing exceeds that of the nominal page-scoring variant; neither statistical significance nor complete-system speedup is established.

Regional and gate diagnostics. A historical protection-only comparison scores 26.09 versus 47.91 with middle-history retrieval, but also reduces the token budget from 1024 to 256; it cannot isolate regional allocation. Gate calibration contains 16 input labels and 36 layers each. Selecting input–layer records whose recorded far-attention weight sum is < 0.05 gives 30 prefill and 22 decode records, with 13 in common. The prefill-selected records have mean decode far-attention weight sum 0.0582 and maximum 0.1678. These stored statistics expose cross-phase changes, not task-quality preservation or request-time savings. A subsequent author report gives five-task means 47.876 ($\epsilon = .05$) and 47.692 (.10), against a reused 47.910 reference: changes of -0.034 and -0.218 points. Each configuration declines on three tasks, with HotpotQA changes $-1.05/-1.26$. Reported five-task bootstrap intervals $[-.584, .452]/[-.806, .400]$ are descriptive, not equivalence evidence. The available execution records do not establish a matched comparison for the shared-subspace configuration (Supplementary Note S5). Matched three-policy quality and complete cost remain unverified.

Available evidence and missing controls. Matched dense FullKV quality, the full-dimensional sparse control, and preparation-inclusive request comparisons remain missing; historical full-L1/tail-32 runs cannot fill them. At nominal 128K, the author-reported E146 RULER summary (main-text Ref. [10]) identifies Qwen3-8B with YaRN factor 4 and a sparse configuration with $\alpha = .125$, $\beta = .375$, $\gamma = .625$. Recomputing unweighted means from its 11 task scores gives 46.43 for FullKV and 23.98 for the sparse configuration, a difference of -22.45 points and a sparse/FullKV score ratio of 51.6%. This is a negative aggregate result, not verified matched validation of the current compact configuration ($\alpha = .25$, $\beta = .125$, $\gamma = .625$): raw predictions are absent from the inspected repository snapshot; paired input identities, scorer execution and the current compact execution identity were not independently verified. For the ordinary fixed-quota path without a quota override, the inspected formula gives $K_{2,\text{mid}} = 1024 - 128 - 128 = 768$ and $\text{round}(128 \cdot .375) \cdot 64 \cdot .625 = 1920$, hence a nominal far-token quota $\max(0, 768 - \min(1920, 768)) = 0$. This inference combines inspected code with reported configuration; the actual execution path has not been independently verified. It does not establish the sole cause of the score gap. The subsequently available 64K files contain aggregate scores and a manifest, not raw predictions; their receipt summaries omit fields including α, β and sink protection, so they do not close full run-configuration identity. No 64K scores are added to this comparison. These summaries do not isolate a causal effect of YaRN or far-region candidate starvation. The historical static-layer policy did not meet its improvement threshold. That negative result and the mixed-policy gate diagnostics constrain interpretation but do not directly test this candidate. Their scores and provenance remain in the evidence register. These Qwen3 diagnostics establish no cross-family or production benefit.

Supplementary Note S7: Budgets and configuration identity

Selected prefix/recent configuration. The currently selected candidate uses four prefix tokens and a 128-token recent window, with dense prefill and sparse decode, while implementation defaults remain 128/128. The author-reported evaluation configuration (E149) explicitly selects 4/128; complete prediction and execution-identity records are unavailable. For sufficient disjoint history, protection is 132 and the nominal middle allocation is $1024 - 132 = 892$: near remains 640 and far receives 252. A four-token prefix occupies a partial 64-token page, so regional selection can reselect up to four protected prefix positions. The far quota therefore does not guarantee 252 distinct nonprotected tokens. StreamingLLM motivates small sink protection, but its LongBench comparison also shows that retaining a larger semantic prefix can help; it does not establish four as optimal for Qwen3 or validate a 128-token recent window. Compared with 128/128 at fixed K, this change releases 124 retrieval slots and cannot isolate the effect of fewer sink tokens. The archived five-task representation scores and historical timing/protection diagnostics retain their own settings; they are not relabeled as 4/128 results.

Scarce retained support. The sufficient-capacity example in Fig. 2b assumes enough positive-probability candidates in each region and disjoint protection; partial-page prefix reselection is accounted for separately. In the implementation, L1-masked logits become zero probabilities after softmax. Regional TopK clips its requested count to region width, not to the number of positive-probability candidates. If that support is too small, zero-probability masked positions may enter through ties. Thus page eligibility is the intended support, not an unconditional no-filler guarantee for every short or degenerate input. Actual selected counts and zero-tie behavior require explicit checks under scarce support.

Fair comparison. The sparse-comparator budget is the final unique attention support K, including protected tokens. L1 candidate pages are an internal design choice, not a quantity that must match a competitor. Their construction, scoring, ranking and memory traffic remain part of cost. Model, input IDs, generation limits, prefill/decode treatment, dense-layer policy and task scorer must be matched. Native and controlled-adapter results must be distinguished. Actual token counts after page rounding and selection granularity must accompany nominal K.

Shared-subspace configuration LongBench settings. The five experimental metadata files specify page size 64, L1 cap 128 pages, final cap 1024 tokens, compression ratio 4, tail coordinates, $\alpha = .25$, $\beta = .125$, $\gamma = .625$, near-average/far-minmax page scoring and 4-bit quantize–dequantize values. They specify shared KV-head selection, dense prefill and disabled static layer skip. All five execution records identify producer `35b39bd8e43cf478bb4f0ab64a8457f08d87c864`. For these non-projection

runs, that tree gathers 32 coordinates for L1 but scatters quantize–dequantize values into zero-filled width-128 L2 keys and scores them with full-width queries. Thus 32/32 describes active coordinates, not both physical contraction widths. The later non-projection shared-subspace configuration path gathers width-32 Q/K for L2; the archived quality score does not directly validate that implementation. These implementation paths do not establish runtime execution or finite-precision ranking equivalence. The recorded tree also gives prefix protection $2 \times 64 = 128$, recent protection 128, and sparse decode on all Qwen3 attention layers. These code-derived fields do not certify an unrecorded modified worktree or an execution trace. The normal long-history allocation is 16/112 near/far pages and 640/128 middle tokens, plus 256 protected tokens. The 8192 candidate slots do not imply 8192 final attention tokens, nor bound L2 scoring: ordinary L2 scores the full history before masking. Realized per-step counts were not recorded.

The configured input cap is 31500 for these runs. Task sample counts and maximum generated tokens are: Qasper 200/128; HotpotQA 200/32; MuSiQue 200/32; GovReport 200/512; RepoBench-P 500/64. The cap is not an observed input length. The complete model/tokenizer revision, runtime precision, hardware identity and scorer revision are not established by these experimental metadata; they must be bound for a matched comparison. Document-cluster identity, a matched dense run and a predeclared task-quality tolerance are still required for non-inferiority testing.

Nominal full-dimensional page-scoring variant. Five legacy prediction files have no manifest or execution record. Filename/verdict records identify page 64, L1 cap 128, final cap 1024, compression ratio 4 and the same nominal α, β, γ. Effective input cap, full implementation revision, protection, phase and complete scorer settings are not verified. Current defaults cannot fill these gaps.

Nominal full-dimensional token-scoring variant. Five manifests contain the same listed budget settings and an explicit full L2 flag. Four receipts omit the implementation commit. The MuSiQue receipt identifies `60027fba3672f33744917c4f59960d0f131f8ab0`, but that tree does not implement the recorded L2-mode field. The code actually executed is therefore not established by that commit. The nominal label and score are retained; full-dimensional L2 execution is not certified.

Quest comparator. The official evaluation uses dense multi-token prefill, sparse token-by-token processing, dense first two layers, and page selection per query head. Its 1024-token budget with chunk size 16 differs internally from 128 pages of size 64 in our selector. A controlled comparison aligns final unique K and phase/layer policy, while reporting the page quotas and head-sharing distinction separately. A paper's native configuration and a controlled adapter must not be pooled into one result. No matched Quest or FullKV quality–request-cost result is claimed here.

Other result families. The synthetic 64K selector timings are single-layer, batch-one, single-query summaries with an unresolved sampling protocol (S6); they are not bound to the LongBench requests. Gate summaries identify thresholds but not complete implementation and subspace settings. Protection-only uses 256 tokens versus its reference's 1024. Missing identity, unknown gate/refill settings, and unmeasured complete costs remain unknown rather than being inferred from current defaults.